Inkloom AI – Brand Governance, Before Publish

Round 1 Submission: The Plan

1. Problem Definition and Evidence

Growing organisations publish content through many employees, agencies and AI tools. Brand guidance is scattered across documents, while review is usually manual and happens late. As a result, inconsistent voice, unsupported claims, outdated facts, wrong asset use and accessibility gaps can reach the public before anyone notices. A single content item may pass through several people using different tools and different versions of the rules. Generic AI tools can also produce fluent but unsupported edits.

Risk pattern	Potential consequence	What Inkloom checks
Voice, terminology or messaging drift	Fragmented identity and confusing customer experience	Tone cues, preferred/prohibited terms and approved message rules
Unverified or stale statements	Misleading promises, credibility loss, avoidable escalation	Claims against dated approved facts; explicit “no source found” fallback
Wrong asset use / accessibility omissions	Visual mismatch or content that excludes users	Asset references and basic text/alt-text checks in the staged MVP

Evidence status

The event brief identifies multi-author, multi-tool publishing, incorrect claims, outdated information and inconsistent voice or visuals as the core challenge. This problem statement is our initial evidence. It is **not** a claim that our team has measured how common the problem is; no primary interview, survey or usage data has been collected yet. Frequency and cost are hypotheses until measured.

Planned research (to be carried out and reported only after collection)

Method	What we will do	Decision it informs
Interviews	Speak with creators, brand/communications managers and reviewers about their last real review, tools used, delays and recurring corrections.	Prioritise checks and define the first workflow.
Short survey	Survey accessible target users on role, content volume, error types and current review effort.	Test the pain and set a pilot baseline.
Scenario test	Users submit a draft, inspect cited findings, revise and route for approval; we record completion and confusion.	Improve explanations and workflow.

Research integrity: Demo data will be synthetic. Personas and impact estimates remain hypotheses until validated. Real confidential content will not be used without permission.

2. Target Users, Needs and Organisational Setting

Setting: a growing organisation with a brand/communications owner, multiple content creators and an approval path before external publication. Likely content includes campaign copy, social posts, web pages and documents. The MVP prioritises text and PDF; richer visual analysis is staged and not

claimed as complete.

User	Need	Inkloom response
Creator / marketer	Fast, actionable feedback without losing intent or ownership.	Findings, explanation, cited source and optional edits.
Brand / communications manager	Consistent enforcement, clear exceptions and less repetitive correction.	Review queue, severity, approval actions and recurring-pattern insights.
Authorised reviewer	Traceable decisions and visibility into high-risk claims.	Role-aware actions, comments, evidence links and revision history.

3. Research, Interviews, Observations and Surveys

As stated in Section 1, no primary research has been completed at the time of this plan. The evidence base today is the event problem statement and its expected capabilities. The validation plan above (interviews, survey and scenario tests) will run first in development (Milestone 1), with a documented baseline. Results will be reported honestly after collection and used to adjust check priorities, the workflow and the explanations.

4. Proposed Solution, Value Proposition and Workflow

Solution: a pre-publication brand-governance platform that reviews draft content against approved guidelines, facts and assets, explains each finding with the source involved, and suggests corrections that preserve the creator's original purpose.

Value proposition: move brand knowledge from scattered documents into one controlled checkpoint. Unlike a generic pass/fail score, Inkloom shows which approved source applies, why content may conflict, and where evidence is missing. Creators may accept, reject or edit any suggestion; designated reviewers keep approval authority.

End-to-end workflow

- **Configure** – The owner adds guideline sections, preferred terms, approved facts and asset references with version and date.
- **Submit** – The creator pastes text or uploads a supported PDF and selects the content context.
- **Review** – The engine extracts text, retrieves relevant sources and returns findings, citations, rationale and suggested fixes, or "no source found".
- **Resolve** – The creator edits, accepts/rejects suggestions or comments; revisions are recorded.
- **Approve** – A reviewer resolves high-risk or unresolved items and approves, requests changes or escalates. There is no automatic publishing.
- **Learn** – A dashboard surfaces recurring issue types and unclear rules for the owner to improve.

Pilot success signal: more reviewed drafts have traceable resolution before publication, with fewer repeated preventable corrections, measured against a baseline rather than assumed.

5. Planned Features, Technology Stack, Integrations and Data Sources

Planned features

- Import or define guidelines, approved facts, terminology and asset references with version/date.

- Paste text or upload PDF for review before publication.
- Detection of tone, terminology, messaging, unsupported-claim and basic accessibility issues (e.g. alt-text presence, reading level).
- Explained findings with cited guideline/source and an optional suggested fix; “no source found” fallback.
- Side-by-side findings view with accept/reject; comments; revision history.
- Role-based workflow (creator, manager, reviewer) with approval queue.
- Insights dashboard for recurring issues and unclear guidelines; AI-use disclosure banner.

Layer	Planned responsibility	MVP boundary
Web UI	Upload/paste; findings view; reviewer queue; comments; role switcher; accessible controls.	Seeded fictional data; no production integrations assumed.
API + workflow	Submission lifecycle, review/approval states, comments, roles and version records.	No publish-on-approve automation; explicit human approval.
Review engine	Text/PDF extraction; retrieval of relevant guideline/fact snippets; structured checks; optional model-assisted tone, explanation and suggestion.	Rules and cited retrieval first; uncertain judgments and missing evidence are labelled.
Data store	Guidelines, sources, submissions, findings, decisions and revisions with source/version references.	Synthetic demo data; retention/deletion defined before any real pilot.
Insights	Aggregate recurring categories, unresolved flags and guideline gaps.	Descriptive trends; no employee ranking.

Note: the exact providers, libraries and AI services will be finalised during the build and fully disclosed in the Round 2 submission, including any AI-generated components.

Data sources and integrations

- Organisation-approved guidelines, terminology, fact sheets, claims and authorised assets, with owner/version/date.
- Submitted drafts and human decisions (comments, accepted/rejected suggestions, approval state) as workflow records.
- A fictional brand pack and synthetic “good / problematic” examples for the prototype.
- Document storage, CMS or identity integrations are possible later, only after permission and security review; they are not an MVP dependency.

Key technical decisions

- **Evidence-linked over confident:** every finding attaches the source passage and version; missing evidence means unknown, not proof of correctness.
- **Structured output:** issue type, severity, explanation, source, suggestion and uncertainty label, validated before display.
- **Human in the loop:** suggestions are optional; exceptions and approvals are attributed and versioned.
- **Formats in stages:** reliable text/PDF extraction first; richer asset checks only with clear limitations and tests.

Round 2 demo target: submit a fictional campaign draft; surface a source-backed conflict and an unsupported claim; revise it; complete manager approval with an audit trail.

6. Responsible Design: Privacy, Security and Human Oversight

Concern	Planned safeguard
Confidential materials	Least-privilege access; protected transport/storage in deployment; minimise content; define retention and deletion. No customer content used for model training without explicit consent and policy.
External AI services	Disclose provider/model, data sent, retention and limits. Prefer private/no-training configurations where available; send only necessary excerpts; keep a basic rule-check path.
False positives / misses	Show evidence and uncertainty; test varied examples; allow corrections. The tool is not legal, factual or accessibility certification; high-risk claims go to authorised reviewers.
Bias and accessibility	Test tone rules across dialects and writing styles; avoid demographic inference or staff ranking. Keyboard-accessible interface, readable contrast and clear labels.
Human authority	No silent rewriting, auto-approval or auto-publishing. Original content and decision history are preserved; humans own final approval.

7. Development Plan, Role Allocation and Expected Impact

Dependencies are sequenced early: research and seed data start first; the engine and workflow API define usable endpoints; the frontend integrates against mock data early, then switches to real endpoints.

Role	Owner	Responsibilities	Why it matters
AI Review Engine Lead	M Tejas Navada	Prompting and retrieval over guidelines/facts; structured findings (issue type, explanation, cited guideline, suggested fix); “no source found” fallback; text and PDF extraction; accessibility check logic (alt text, reading level).	Core of the product. Without reliable, cited findings there is nothing to demo; carries most technical depth and explainability.
Backend and Workflow Lead	Aarush C U	Database (guidelines, content, findings, comments, versions); review/approval states; role switcher (creator/manager/reviewer); revision history; API connecting engine to UI.	The end-to-end loop depends on it; approval workflow is a named requirement.
Frontend and UX Lead	K H Varun	Upload screen; side-by-side findings view with accept/reject; reviewer queue; insights dashboard; AI-use disclosure banner; accessible UI.	Judges experience the product through this; built against mock data early, then real endpoints.
Research, Data and Documentation Lead	Karthik R	Survey/interviews; fictional brand seed guidelines and sample bad content; PDF/report; architecture diagram; demo script and video; disclosure of AI tools used.	Lowest build dependency but starts first; makes the demo realistic.

Milestones

Milestone	Exit check
1 – Validate and define	Initial user input; documented assumptions and baseline measures; fictional seed pack.
2 – Engine + workflow	API returns source-linked findings; persists versions and review states; tests missing-source and failure cases.
3 – Usable demo	Creator → review → manager approval end-to-end; role-specific UI; usability feedback and fixes.
4 – Evidence and handoff	Scenario tests; limitations, architecture, dependencies, privacy choices and contributions documented for Round 2.

Expected impact

Reduce avoidable review friction, identify off-brand or unsupported content earlier, and turn repeated corrections into clearer guidance. Impact will be evaluated with pilot measures (review completion, traceable resolution, repeated issue rate and user-reported usefulness) rather than by claiming unmeasured gains.

8. Coverage of the Problem Statement Capabilities

The table maps each capability named in the problem statement to the plan, and states what is delivered in the MVP and what is staged for later, so that scope stays honest and feasible.

Capability required	How Inkloom addresses it	MVP / Staged
Import or define guidelines, facts, terminology, assets	Owner enters guideline sections, preferred/prohibited terms, dated approved facts and asset references with version and date.	MVP
Review text, images, documents or other formats	Paste text and upload PDF with text extraction first; image and richer asset checks added once tested.	Text/PDF: MVP. Images: staged
Detect tone, language, facts, claims, visual usage, accessibility	Rule and retrieval checks for terminology, claims and alt-text presence; model-assisted tone judgments labelled with uncertainty. Visual-usage checks use asset references.	Text checks: MVP. Visual: staged
Explain the flag and identify the source	Each finding shows issue type, severity, rationale and the cited guideline passage with version; “no source found” when evidence is missing.	MVP
Suggest corrections without overriding humans	Suggestions preserve the creator’s intent and are optional; original text is kept and nothing is changed silently.	MVP
Comments, approvals, revision history, roles	Creator, manager and reviewer roles; comments on findings; approve, request changes or escalate; every revision and decision recorded.	MVP
Recurring problems and governance insights	Dashboard of repeated issue types, unresolved flags and guideline sections that cause confusion.	MVP (basic)
Protect confidential material; disclose AI use	Least-privilege access, minimal data sent to external services, no training on customer content, visible AI-use disclosure banner.	MVP

Illustrative finding (fictional example)

To show what “explainable” means in practice, a finding in the prototype is planned to look like this:

Field	Example
Draft sentence	“Our app is the fastest in the world and guarantees zero errors.”
Issue type / severity	Unsupported promotional claim / High
Explanation	The sentence makes an absolute performance claim. The approved fact sheet does not contain a verified claim of this kind.
Cited source	Fictional Brand Pack, Claims Policy, section 2.1 (version and date shown)
Suggested fix (optional)	“Our app is built for speed and designed to help you reduce errors.”
Uncertainty label	Source found / rule-based. Routed to an authorised reviewer because severity is High.

MVP scope

- **In scope:** text and PDF review, cited findings, accept/reject, comments, role-based approval, revision history, insights dashboard, synthetic data.
- **Out of scope for MVP:** automatic publishing, real customer data, CMS or identity integrations, full image/video analysis, legal or factual certification.

9. Risks, Testing, Success Measures and Disclosure

Risk	Mitigation
Model gives fluent but wrong findings	Cited retrieval first; every finding needs a source or an explicit “no source found” label; structured output validated before display.
Too many false alarms frustrate creators	Severity levels; creators can reject with a reason; rejection patterns feed the insights dashboard to improve unclear rules.
Scope too large for the build window	Text/PDF first; visual checks staged; frontend built on mock data early; milestones have clear exit checks.
External AI service unavailable or restricted	Basic rule-check path keeps the product working without the model; only necessary excerpts are sent.
Confidential content exposed	Synthetic data only in the prototype; retention and deletion defined before any real pilot.

Testing and reliability plan

- Test set of fictional “good” and “problematic” drafts covering tone, terminology, stale facts and unsupported claims.
- Failure cases: missing source, empty or unreadable PDF, conflicting guidelines, service timeout.
- Workflow tests for each role: creator edits, manager approval, reviewer escalation, revision history.
- Usability check with a few target users, with problems found recorded and fixed.

Success measures

Measure	How it is measured
Review completion	Share of submitted drafts that reach an approve, request-changes or escalate decision.
Traceable resolution	Share of findings with a recorded accept, reject or edit and a linked source.
Repeated issue rate	How often the same issue type recurs per content type over time, compared with the baseline.
User-reported usefulness	Short post-test rating and comments from creators, managers and reviewers.

Disclosure and event-rule compliance

- All development will take place within the official build window; work before it is limited to research and planning, as allowed.
- AI tools, libraries, templates, datasets and external services will be disclosed, including significant AI-generated components.
- Open-source components follow their licences; demo data is fictional. Each member’s contribution will be stated accurately.

Future roadmap (after the MVP)

- Image and visual-consistency checks against authorised asset libraries.
- Document storage, CMS and identity integrations after security review; richer insights and multi-language support.